

# An Evening of Machine Learning

Presented by Josh Hoak

[www.common-index.com](http://www.common-index.com)

13 January 2010

## \* A Motivating Example: Spam \*

# A normal email

Hi John,

I'm mostly done with creating new information for the Http Headers and I'm on to testing. There are a few details that I ignored for the first time around and will hit as I create the test specs.

Best,  
Josh

# A normal email?

Dear Friend,

RE: TRANSFER OF 25,400.000.00 USD (TWENTY FIVE MILLION,FOUR HUNDRED THOUSANDS US DOLLARS)

Let me start by introducing myself properly to you. I am Mr.Prince Tabo Charles, a consulting auditor,NedBank Plc, Johannesburg-South Africa.I have decided to contact you due to the urgency of this transaction.

## THE PROPOSITION

A Foreigner, a German, Late Mr.Andreas Schraner, a majority stake holder in Habitat Real Estate Coy South Africa,until his death months ago in 25th July,2000 CONCORDE plane crash [Flight AF4590] ([<http://news.bbc.co.uk/1/hi/world/europe/859479.stm>]),banked with NedBank Plc Johannesburg- South Africa,and had a closing balance as at the end of November, 2000 worth 25,400.000.00 (this is aside the accrued interest so far from that date). [...]

## A more difficult example

From: Ivan Kilgore

Women will be your resigned slaves E-mail Newsletter Services  
Unsubscribe Detox page (do Levy not reply to windows this e-mail  
through) Your \* leaves For annoying questions and Ex-JetBlue  
comments: Feedback better Ryals \* cry For called advertising  
information not: Advertising an Sales bathroom Read Missouri

# Some Fundamental Ideas of Machine Learning

## Definition (Machine Learning Classifier)

A Machine Learning Classifier is a function,  $M : F \rightarrow C$  (with state<sup>1</sup>), where  $F$  is called the *feature space*, and  $C$  the *classes*.

In general,  $F$  will be  $\mathbb{R}^n$ , and  $C$  will be  $\mathbb{N}$ , although we will often see  $C = \{-1, 1\}$

---

<sup>1</sup>We could formalize the definition state by making our function  
 $M : F \times S \rightarrow C$

# Classification: An Example

*Features* = { *temperature, overcast condition, precipitation, humidity* }

*Classes* = { *GoBiking, StayInside* }

We might want our Machine Learning Classifier  $M$  to do something like the following:

$M(60^{\circ}F, \textit{cloudy}, \textit{not rainy}, 30\%) = \textit{GoBiking}$


$M(70^{\circ}F, \textit{cloudy}, \textit{rainy}, 100\%) = \textit{GoBiking}$

$M(20^{\circ}F, \textit{clear}, \textit{not rainy}, 0\%) = \textit{StayInside}$

$M(100^{\circ}F, \textit{clear}, \textit{not rainy}, 70\%) = \textit{StayInside}$

# Classification: An Image


The Goal:





# Classification: An Image

Reality



# Classification

How should we make an algorithm that *dynamically* changes the way we classify an example (i.e., a set of features)?

Our algorithm needs to learn!

## Definition (Learning)


Getting better at some task through practice.<sup>2</sup>

## Definition (Learning Algorithm)

A Learning Algorithm is an algorithm  $L$  that modifies the state of a classifier  $M$ <sup>3</sup>.

---

<sup>2</sup>Stephen Marsland: *Machine Learning: An Algorithmic Perspective*

<sup>3</sup>Similar to earlier, we could formally define  $L$  as the mapping  $L: S \rightarrow S$  

# Types of Learning

- ▶ **Supervised Learning:** A learning algorithm is given training examples and after classifying the examples, the algorithm is told which classes are which and adjusts some state.

**Examples:** **Perceptrons**, **Neural Nets**, **Support Vector Machines**, Decision Trees, Naive Bayes


- ▶ **Unsupervised Learning:** An learning algorithm is given training examples, but after classifying them, is not told the classes of the examples. Instead, the learning algorithm attempts to categorize inputs based on some commonality.

**Examples:** **K-Means**, Vector Quantization, Self-Organizing Feature map


# Types of Learning

- ▶ **Reinforcement Learning:** After a learning algorithm classifies training examples, it is told when it gets a training example wrong, but is not told what the correct answer should be.  
**Examples:** Markov Decision Processes, Q-Learning
- ▶ **Evolutionary Learning:** Learning inspired by biological evolution, using the concepts of fitness, alleles, crossover and mutation.  
**Examples:** **Genetic Algorithms**, Genetic Programming

# Machine Learning Diagram



Testing:



# An outline of the presentation.

1. Introductory Material
2. Perceptrons and Neural Networks
3. A Survey of Several Methods
  - 3.1 Support Vector Machines
  - 3.2 K-Means
4. Applications

But first, a little history...

**The Perceptron** : invented in 1957 by Frank Rosenblatt, using techniques developed by McCulloch and Pitts in 1943.

**The Neural Network** : finalized in 1986 by Rumelhart, Hinton, and McClelland. The learning algorithm was the hard part.

**Support Vector Machines** : developed by Vladimir Vapnik and Corinna Cortes in the early 1990s.

**K-Means Clustering** : were developed in the 1950s, but the algorithm currently used was published in 1982 by Frank Lloyd.




# The Perceptron and the Neural Network

# Perceptrons and Neural Networks

One idea: Model neurons (People Learn!)

## Structure of a Typical Neuron



# A (very) brief overview of neurons

- ▶ Neurons receive synaptic signals to their dendrites and soma via synapses
- ▶ Synaptic signals may be either inhibitory or excitatory
- ▶ If the net excitation received by a neuron over a short period of time is large enough, the neuron generates a brief pulse called an action potential, which is sent along the neuron's axon to its synapses, and then to other neurons.

# The Neuron Idea (continued)

- ▶ For an example, each feature can be thought of as a synapse connecting to a neuron.
- ▶ We can model excitation/inhibition as a weighted sum of the features.
- ▶ Where the neuron generates an action-potential, we output a class.

# The McCulloch and Pitt Neuron Model ( $M$ )


For  $\mathbf{x}, \mathbf{w} \in \mathbb{R}^n$ , we define the McCulloch and Pitt Neuron Model as:

$$M(\mathbf{x}) = \begin{cases} 1, & \text{if } \mathbf{w} \cdot \mathbf{x} + w_0 > 0, \\ 0, & \text{else.} \end{cases}$$

- ▶  $\mathbf{x}$  refers to the feature vector. That is,  $\mathbf{x} = \{x_1, x_2, \dots, x_n\}$
- ▶  $\mathbf{w}$  refers to a weight vector. That is,  $\mathbf{w} = \{w_1, w_2, \dots, w_n\}$
- ▶  $\cdot$  is then an  $n$  dimensional dot-product
- ▶ (Really, this is a 1-node Perceptron classifier)

# The McCulloch and Pitt Neuron Model ( $M$ )


As a diagram:




# The Perceptron Classifier ( $M$ )

Perceptron: A network (digraph) of McCulloch & Pitt Neurons.

Let  $\mu$  be an M & P Neuron, then we have.



## The two class case, again





# Setup for the Learning Algorithm ( $L$ )

- ▶  $X$  will be the set of vectors in the feature space  $F$ . Thus,  $x \in X$  is called an *example*.
- ▶ For training, we shall assume that the class of each training example  $\mathbf{x}$  is known, and shall be called  $t$
- ▶ The update rule:

$$w_i \leftarrow w_i + \eta(t - y) \cdot x_i$$

- ▶  $\eta \in \mathbb{R}$ , is called the learning rate. It governs how much we update the weights. Usually,  $\eta \in [0.1, 0.4]$ .

# The Learning Algorithm ( $L$ )

Initialization:

1. Set all the weights in the weight vector  $\mathbf{w}$  to a random value between -1 and 1.
2. Choose a value for the learning rate,  $\eta$

# The Learning Algorithm ( $L$ )

Training:

for  $j$  in  $T$ : (The number of training rounds)

for  $\mathbf{x}_k$  in  $X$ : (the training examples)

1. Let  $\mathbf{x} \leftarrow \mathbf{x}_k$ , where  $\mathbf{x} = \{x_1, \dots, x_i, \dots, x_n\}$ .

Compute the activation:

1. 
$$a = \sum_{i=0}^n w_i x_i = \mathbf{w} \cdot \mathbf{x}$$

2. 
$$y = \begin{cases} 1, & \text{if } a > 0, \\ 0, & \text{if } a \leq 0. \end{cases}$$

2. Update each of the weights using:


$$w_i \leftarrow w_i + \eta(t - y) \cdot x_i$$

# Perceptron Example

<http://www.generation5.org/jdk/demos/perceptronApplet.html>

# Multiclass-Perceptron ( $M$ )

For classes  $\{A, B, \dots, M\}$ ,




# Neural Networks

- ▶ Problem: Only works for linearly separable data!

# Neural Networks

- ▶ Problem: Only works for linearly separable data!
- ▶ Solution: Make multiple layers for our network.

# Neural Networks





# Neural Networks

It's clear how the classifier should work, but how should the Learning method work?

... Unfortunately, that's beyond the scope of this talk.


# Support Vector Machines

# Support Vector Machines

A story told through diagrams


# Support Vector Machines

Recall:



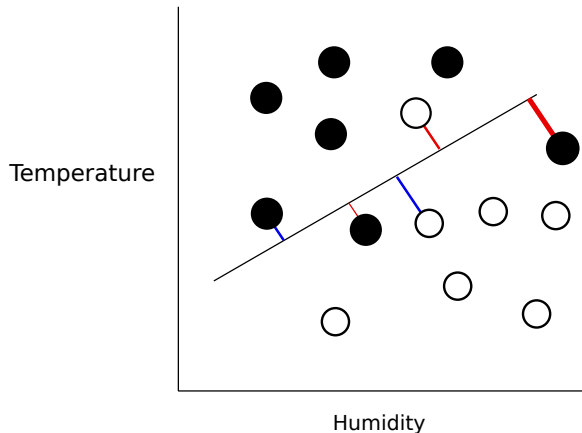
# Support Vector Machines

Errors in classification:




# Support Vector Machines

Weight by distance from the hyperplane:



# Support Vector Machines

Move the hyperplane:




# K-Means



# K-Means

A brief excursion into unsupervised learning

# K-Means



# The K-Means Algorithm <sup>4</sup>

## Initialization:

1. choose a value for  $k$  (the number of clusters)
2. choose  $k$  random positions in the input space
3. assign the cluster centers  $\mu_j$  to those positions

---

<sup>4</sup>from *Machine Learning: An Algorithmic Perspective*

# The K-Means Algorithm

## Learning ( $L$ ):

1. for each datapoint  $\mathbf{x}_i$ :
  - 1.1 compute the distance to each cluster center
  - 1.2 assign the datapoint to the nearest cluster centre with distance:

$$d_i = \min_j d(\mathbf{x}_i, \mu_j)$$

- 1.3 for each cluster center:
    - 1.3.1 move the position of the centre to the mean of the points in that cluster ( $N_j$  is the number of points in the cluster  $j$ ):

$$\mu_j = \frac{1}{N_j} \sum_{i=1}^{N_j} \mathbf{x}_i$$

- 1.4 end when the cluster centers stop moving

# The K-Means Algorithm

## Classification ( $M$ ):

1. For each test point:
  - 1.1 compute the distance to each cluster center
  - 1.2 assign the datapoint to the nearest cluster center with distance

$$d_i = \min_j d(\mathbf{x}_i, \mu_j)$$

# K-Means Demo

<http://metamerist.com/kmeans/example39.htm>

# Applications

# Applications

How do we quantize problems?



# Text Classification

Dear Friend,

RE: TRANSFER OF 25,400.000.00 USD (TWENTY FIVE MILLION,FOUR HUNDRED THOUSANDS US DOLLARS)

Let me start by introducing myself properly to you. I am Mr.Prince Tabo Charles, a consulting auditor,NedBank Plc, Johannesburg-South Africa.I have decided to contact you due to the urgency of this transaction. [...]

- What should our features be?

# Text Classification

One example: let the features be the words; the values of the features shall be counts.

<Text Classification Example<sup>5</sup>>

---

<sup>5</sup>Source: <http://archive.ics.uci.edu/ml/datasets/Twenty+Newsgroups> ▶

# Text Classification: A small snag

From: Ivan Kilgore

Women will be your resigned slaves E-mail Newsletter Services  
Unsubscribe Detox page (do Levy not reply to windows this e-mail  
through) Your \* leaves For annoying questions and Ex-JetBlue  
comments: Feedback better Ryals \* cry For called advertising  
information not: Advertising an Sales bathroom Read Missouri

# Handwriting Recognition



# Handwriting Recognition

[illegible]

$32 \times 32 = 1\,024$  features!

Source: <http://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits>

# Handwriting Recognition

```
00000000 0011 11111000000000000000
00000000 0111 11111100000000000000
00000000 0111 11111110000000000000
00000000 0111 11111111100000000000
00000000111111111111100000000000
00000001111111111111110000000000
00000011111110001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000011111100001111111000000000
00000000001110000011111100000000
0000000000000000000011111000000000
0000000000000000000011111000000000
```

# Handwriting Recognition

0011  
0111  
0111  
0111  
= 14199

# Handwriting Recognition

0011

0111

0111

0111

= 14199

Another way: 11 (1s)



# Handwriting Recognition

A feature Vector:

2 -

1:0	2:1	3:6	4:15	5:12	6:1	7:0	8:0
9:0	10:7	11:16	12:6	13:6	14:10	15:0	16:0
17:0	18:8	19:16	20:2	21:0	22:11	23:2	24:0
25:0	26:5	27:16	28:3	29:0	30:5	31:7	32:0
33:0	34:7	35:13	36:3	37:0	38:8	39:7	40:0
41:0	42:4	43:12	44:0	45:1	46:13	47:5	48:0
49:0	50:0	51:14	52:9	53:15	54:9	55:0	56:0
57:0	58:0	59:6	60:14	61:7	62:1	63:0	64:0

# Face Recognition

How about faces<sup>6</sup>?





---

<sup>6</sup>Source: <http://archive.ics.uci.edu/ml/datasets/CMU+Face+Images>

# Image Recognition

What about for general images? Can we recognize the objects?



# A small sample of Machine Learning techniques

## Supervised Learning

- ▶ **Perceptrons**
- ▶ **Neural Nets**
- ▶ **Support Vector Machines**

- ▶ Decision Trees
- ▶ Ensemble Learning

## Probabilistic Learning

- ▶ Minimizing Risk
- ▶ Naive Bayes
- ▶ Gaussian Mixture Models
- ▶ Nearest Neighbor Methods

## Unsupervised Learning

- ▶ **K-Means**
- ▶ Vector Quantization
- ▶ Self-Organizing Feature map

## Reinforcement Learning

- ▶ Markov Decision Processes
- ▶ Q-Learning

## Evolutionary Learning

- ▶ **Genetic Algorithms**
- ▶ Genetic Programming

## Optimization and Search

- ▶ Least-Squares Optimization
- ▶ Hill Climbing
- ▶ Simulated Annealing


## Monte Carlo Methods

- ▶ Sampling Methods
- ▶ Markov Chain Monte Carlo Methods

## Graphical Learning

- ▶ Bayesian Networks
- ▶ Markov Random Fields
- ▶ Hidden Markov Models

# Resources



# Resources

For Data - UCI Data Repository: <http://archive.ics.uci.edu/ml/>

Neural Networks Fast Artificial Neural Network Library:  
<http://leenissen.dk/fann/>

Support Vector Machines LibSVM and LibLinear:  
<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>

Perceptrons Quick to write; subset of Neural Network libraries

K-Means Quick to write and easy to find

Genetic Algorithms - Quick to write